

Pending Tables Preview (cutoff-60 protocol)

draft

draft

1 Paired Fair Baseline

Table 1. Paired comparison against an equal-capacity zero-shot LLM on the same 500-claim probe set (mild perturbation). Both sides are binarized with the identical rule: REAL iff score ≥ 60 . ΔAcc = paired accuracy difference; CI = 10,000-sample bootstrap 95% CI; McNemar exact two-sided test on discordant pairs.

System	Acc.	Prec.	Rec.	F1	TP/FP/FN/TN
NewsAnalyzer (hybrid)	99.0%	99.1%	99.8%	0.994	442/4/1/51
Zero-shot Qwen3.8-27B	88.2%	94.9%	91.7%	0.932	406/22/37/33
$\Delta\text{Acc} = +10.85$ pp, 95% CI [8.03, 13.65], McNemar $p = 1.2 \times 10^{-14}$ ($n = 498$)					

Table 2. Component ablation under the fixed decision rule ($n = 500$, offline recomputation from collected per-signal scores; ΔAcc relative to the full system). “Evidence” = fact-check census + community signal; “Metadata” = domain, sentiment, similarity, timeliness.

Configuration	Acc.	ΔAcc	Prec.	Rec.
Full system	98.6%	—	98.7%	99.8%
Evidence-only	98.8%	+0.2 pp	98.9%	99.8%
Metadata-only	98.2%	−0.4 pp	99.3%	98.7%
w/o domain reputation	98.6%	± 0.0 pp	98.7%	99.8%
w/o sentiment volatility	98.8%	+0.2 pp	98.9%	99.8%
w/o temporal decay	98.6%	± 0.0 pp	98.7%	99.8%
w/o safety clamping	95.8%	−2.8 pp	98.6%	96.6%

Table 3. Safety clamping (Stage 3) on/off under the fixed decision rule, stratified dev/test split (seed 42). With clamping enabled, all $\tau_{\text{high}} \in [25, 50]$ yield identical decisions; $\tau_{\text{high}} = 25$ is the deployed configuration.

Configuration	Dev ($n = 250$)		Test ($n = 250$)	
	Acc.	TP/FP/FN/TN	Acc.	TP/FP/FN/TN
No clamping (Stage 3 off)	95.2%	212/2/10/26	97.2%	217/2/5/26
Full clamp ($\tau_{\text{high}}=25$)	99.2%	222/2/0/26	98.8%	221/2/1/26

Table 4. Empirical fake rate per final-score band ($n = 500$; sample composition 444 FAKE / 56 REAL). Bands correspond to the deployed rating vocabulary.

Score band	Rating	n	$\mathbf{P}(\text{FAKE} \mid \text{band})$
0–20	Highly suspicious	282	.993
20–40	Suspected fake	118	.983
40–60	Unverified	46	1.000
60–75	Likely credible	3	.333
75–100	Highly credible	50	.000